


# Web Based Graphics & Virtual Reality Systems

## Programmable Shaders






# CPU and GPU


- The concept of GPU starts from early 2000
- Specially tailored for processing rendering tasks of graphics
- GPU become programmable in its rendering process
  - Shader will be introduced later in this lecture

# CPU and GPU

- Basic difference between CPU and GPU
  - Large number of parallel units in the pipeline
- This design is especially useful for graphics applications
  - Large number of vertices undergo similar or same transformation and operations
- Major GPU manufacturers
  - Nvidia : Geforce Series
  - AMD / ATI : Radeon Series
  - Intel



# Computing Power of CPU and GPU



**GFLOP :**  
Floating-point  
Operations per  
Second



# Computing Power of CPU and GPU

- Compute

- Intel Core i7 – 4 cores – 100 GFLOP
- NVIDIA GTX280 – 240 cores – 1 TFLOP

- Memory Bandwidth

- System Memory – 60 GB/s
- NVIDIA GTX680 – 200 GB/s
- PCI-E Gen3 – 8 GB/s

- 2012 – GPU peak FLOP about 10x CPU

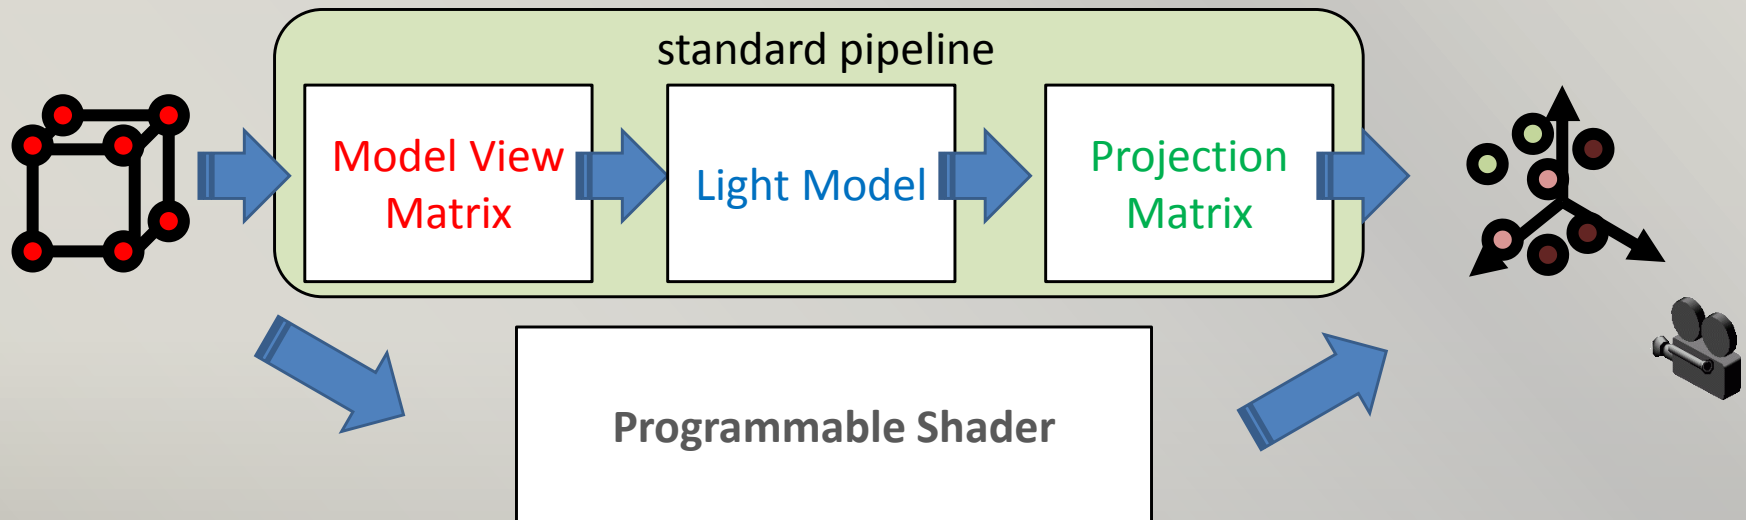
# Programmable Shader

- The standard rendering pipeline are being hard-wired in graphics hardware for years
- Until around 10 years before, the graphics hardware becomes programmable
- The small programs which can customize certain parts of the rendering pipeline are called “Shader”



# Programmable Shader

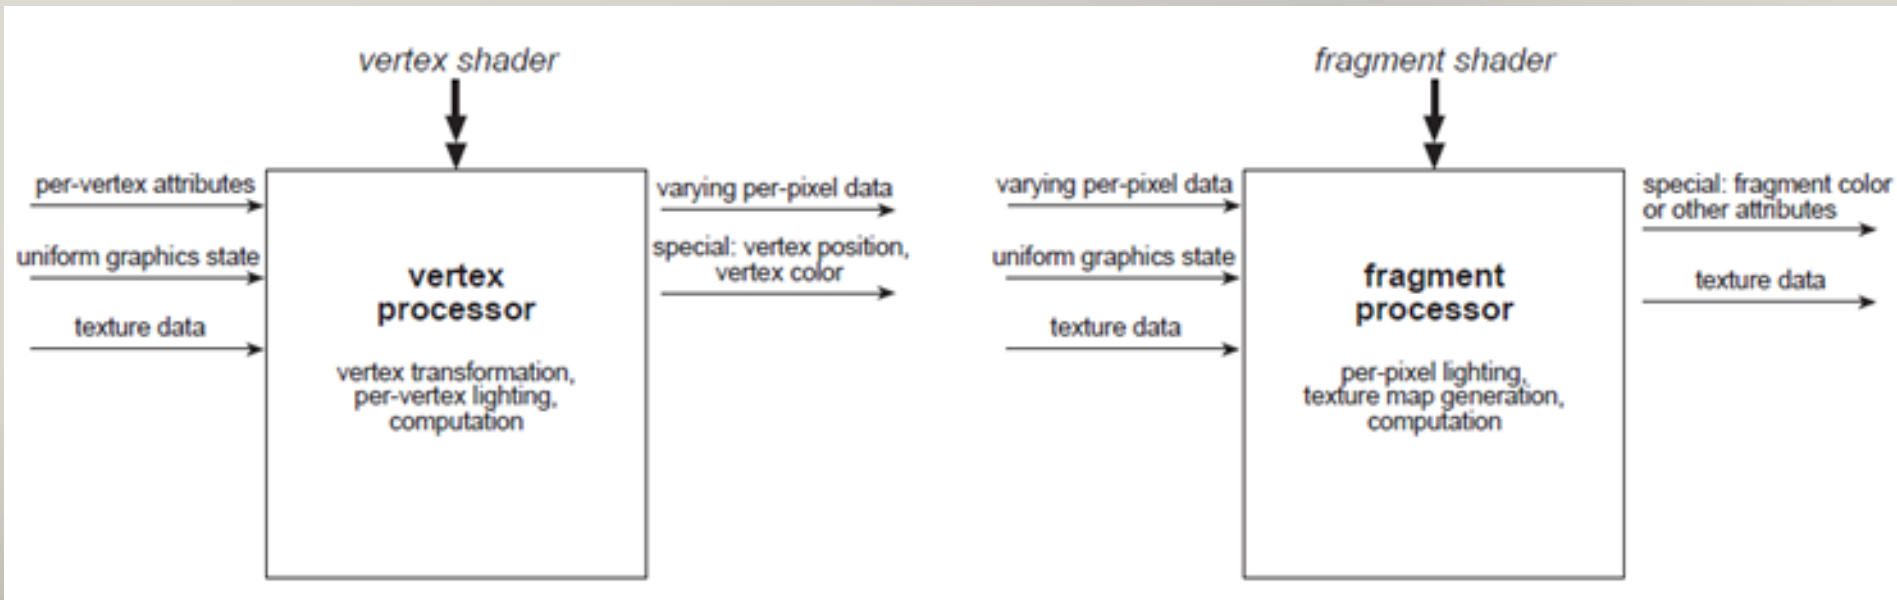
- Instead of going through standard pipeline
- We can now write own programs to perform other processing steps wanted
- E.g. for vertex processing :





# Programmable Shader

- Two major kinds of shaders are available:
  - Vertex shader (for vertex processing)
  - Fragment shader (for fragment processing)





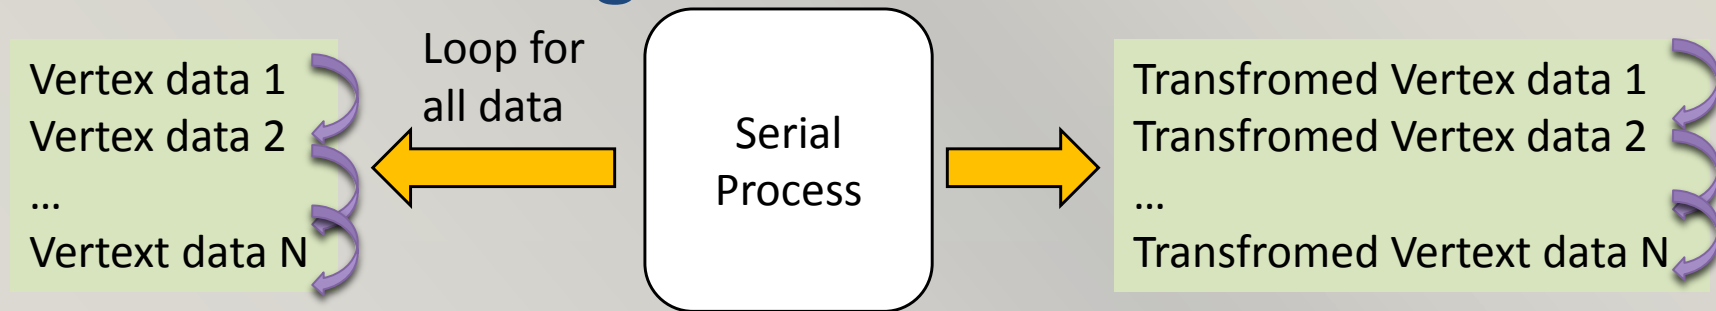


# Programmable Shader

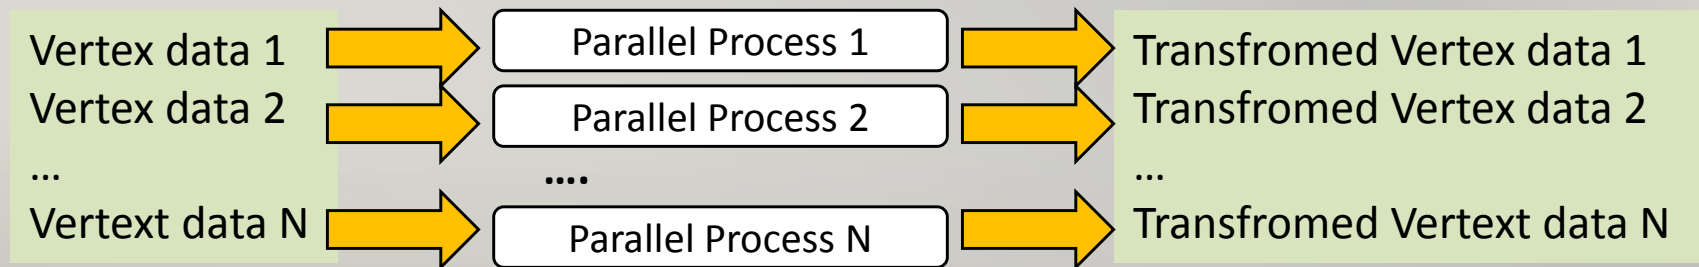
- ✕ GPU hardwares are designed to process a large number of vertices and fragments at the same time
- The major difference of GPU shaders to normal programs is that they are being processed in a parallel manner
- Serial program process all data one by one
- Parallel programs processes all data at the same time

# Programmable Shader

## ■ Serial Processing

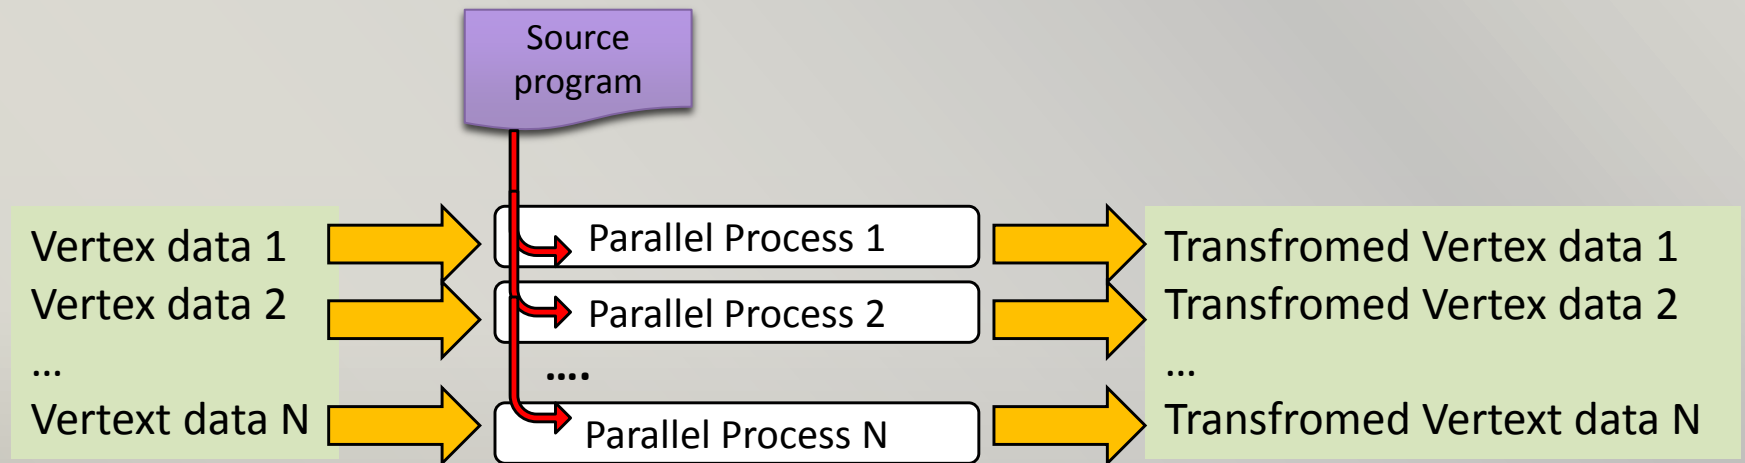


## ■ Parallel Processing




# SIMD

- The parallel processing model in GPU is called SIMD (Single Instruction Multiple Data )
- SIMD architecture
  - Incoming data are different , but all program source are the same



# Programmable Shader

- Although shader is run inside GPU, several languages are available for writing it
  - Cg (Nvidia)
  - GLSL (Open standard)
  - CUDA (Nvidia)
  - OpenCL (Open standard)






# Shaders

- Various kinds of shaders are available
  - Vertex Shader
  - Fragment Shader
  - Geometry Shader
- Responsible for handling different elements in corresponding part of the rendering pipeline

# Applications of Shaders

- Special and advanced real-time effects can be done with the use of shaders




(d)  
Bump mapping

Volumetric lighting



Shadow mapping




Skin rendering



# Application: Physics Simulation


- PhysX engine (Nvidia)
- Rigid motions
- Cloth simulation
- Fluid dynamics





# Application on Rendering Approaches

- Real-time non-photorealistic rendering
- Real-time Raytracing




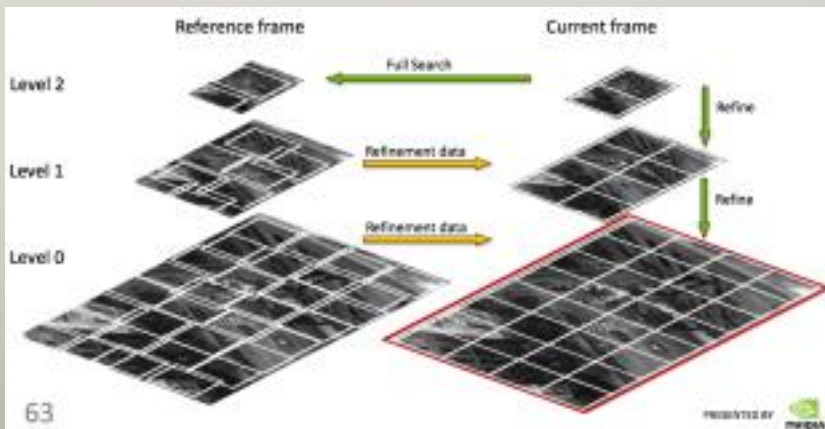

Nvidia GPU Raytracing



Toon Shaded Game

# Application: Image and Video Processing

- Image Registration
- Motion Estimation
- Video Encoding and Compression
- Mainly to accelerate computational intensive tasks




# GPU Architecture: the basics

- CPU runs machine or assembly code of your compiled program
- GPU runs also assembly code, but from your compiled shader

Shader

```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```

Compile

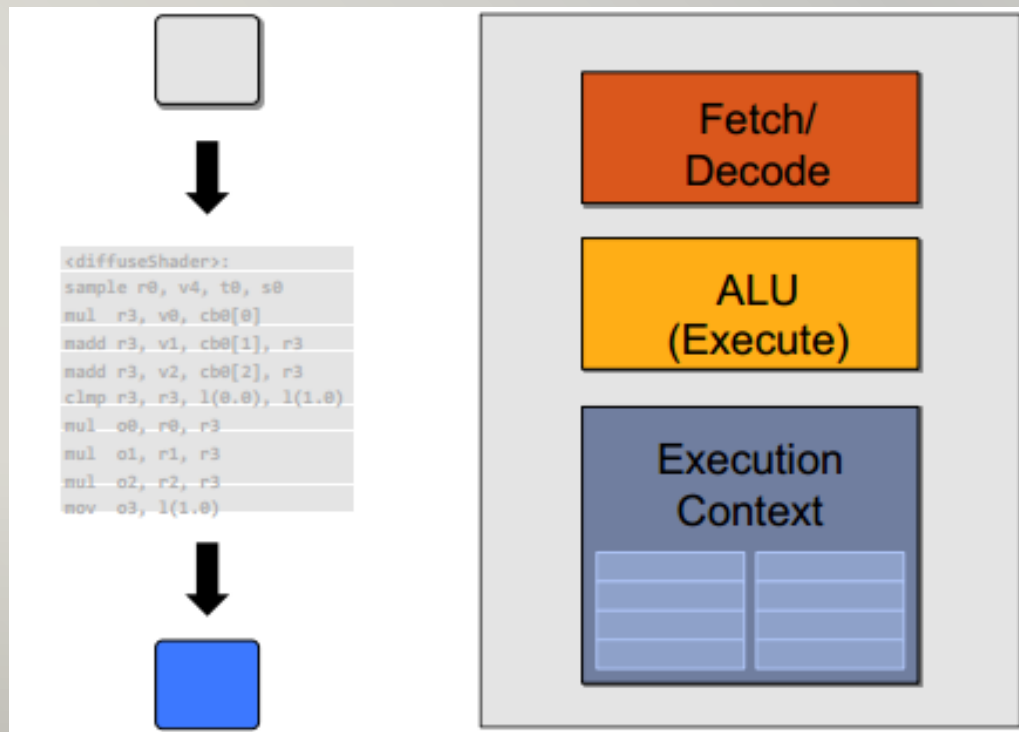


Assembly code

```
<diffuseShader>  
sample r0, v4, t0, s0  
mul   r3, v0, cb0[0]  
madd  r3, v1, cb0[1], r3  
madd  r3, v2, cb0[2], r3  
clmp  r3, r3, l(0.0), l(1.0)  
mul   o0, r0, r3  
mul   o1, r1, r3  
mul   o2, r2, r3  
mov   o3, l(1.0)
```

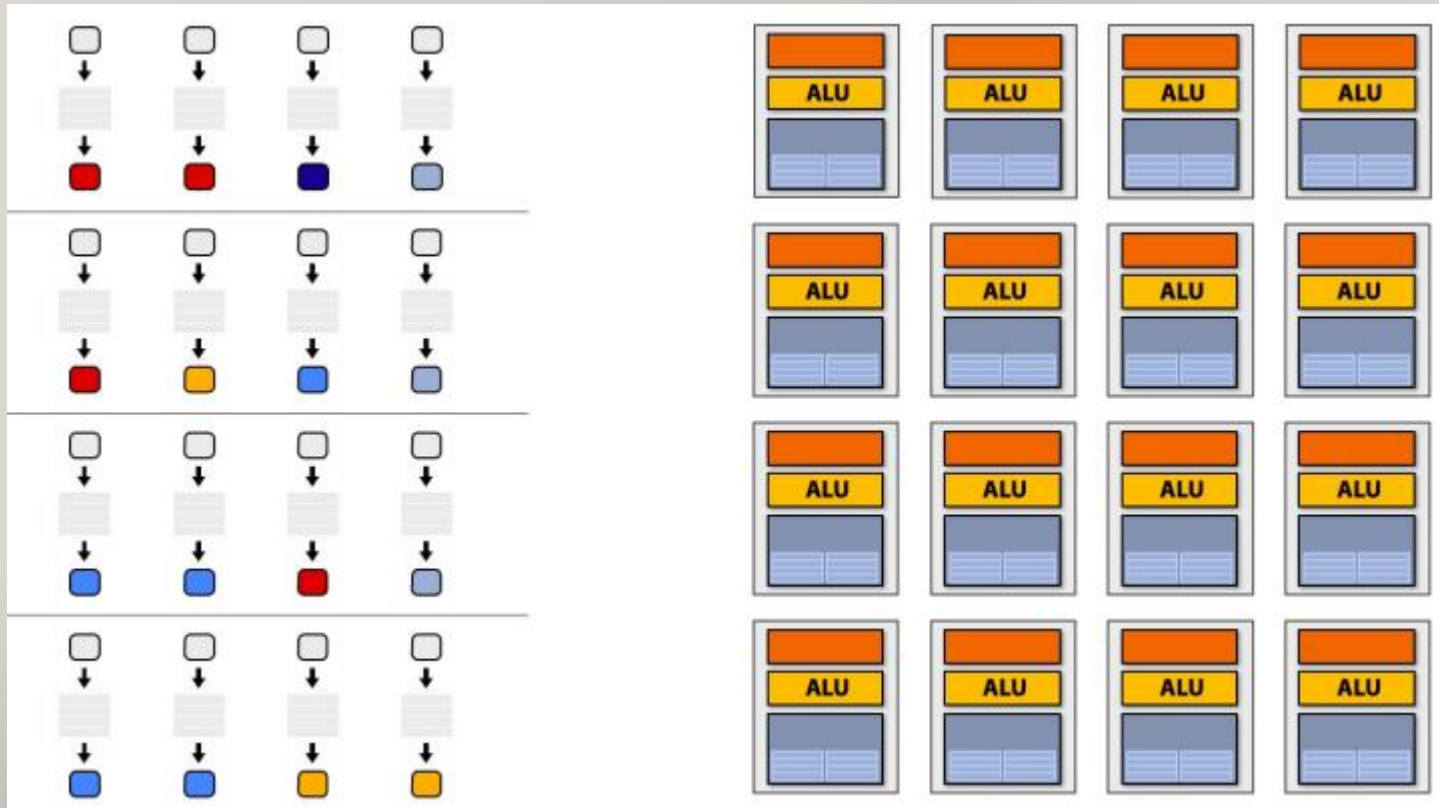
# Cores in GPU

- Each GPU core runs the shader for a particular element (e.g fragment)



# Cores in GPU

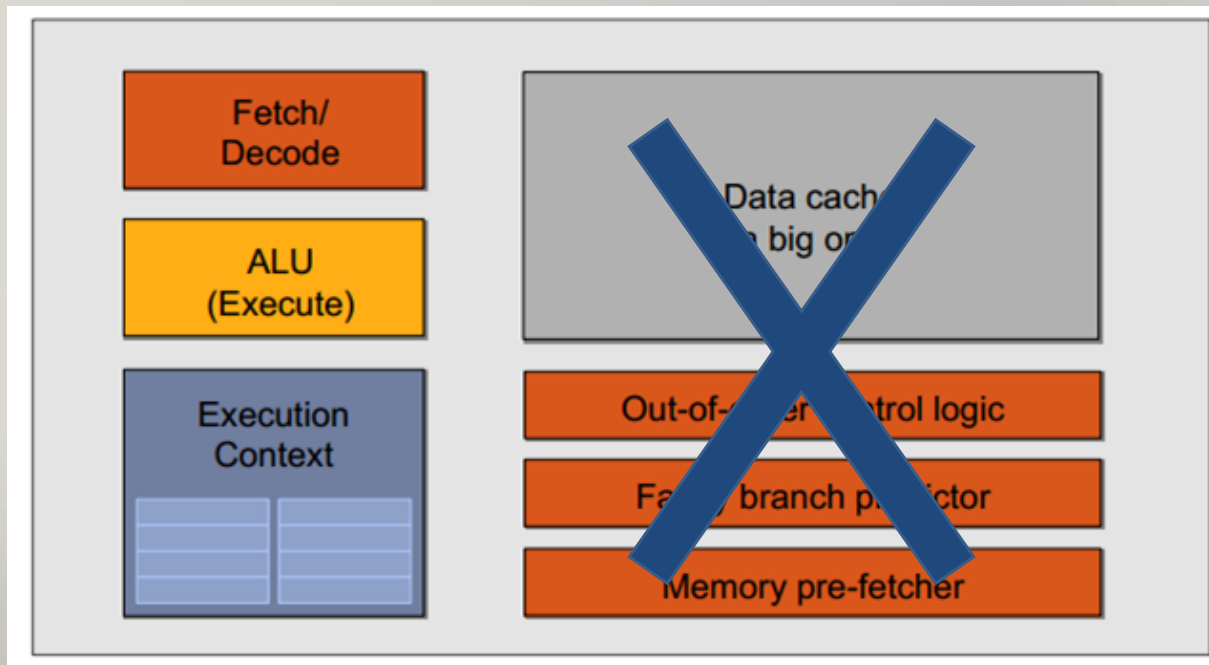
- Multiple cores running in parallel





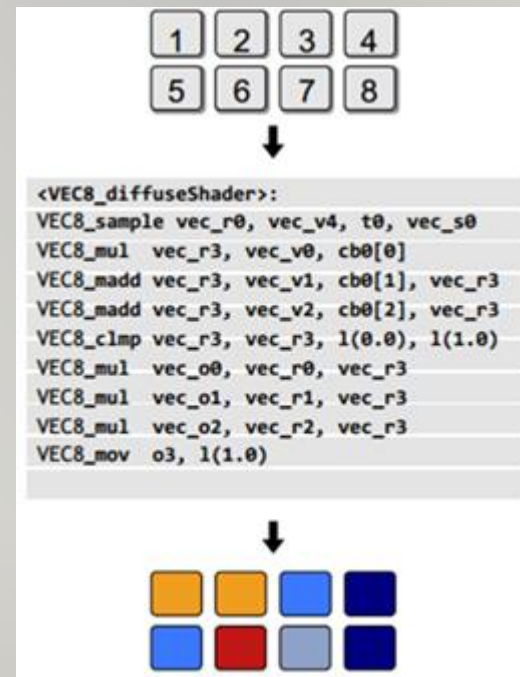
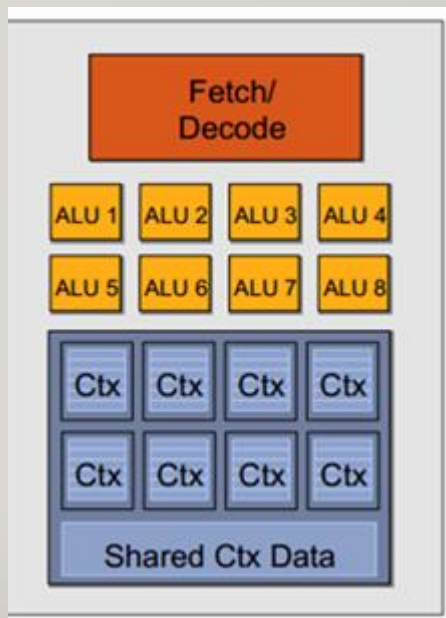
# Cores in GPU

- Compare to a CPU-style core, GPU keeps a single core as simple as possible
  - Referred as “Slimmed down core”



# Cores in GPU

- Tailor for SIMD processing
  - Combine fetch unit of several elements in one core






# Cores in GPU

- In practice, 16 to 64 fragments share an instruction stream




- Multiple cores runs in a GPU
  - Each of them can run same or different shader program




# Cores in GPU

- Each core processes several elements simultaneously
  - Several ALU in one core



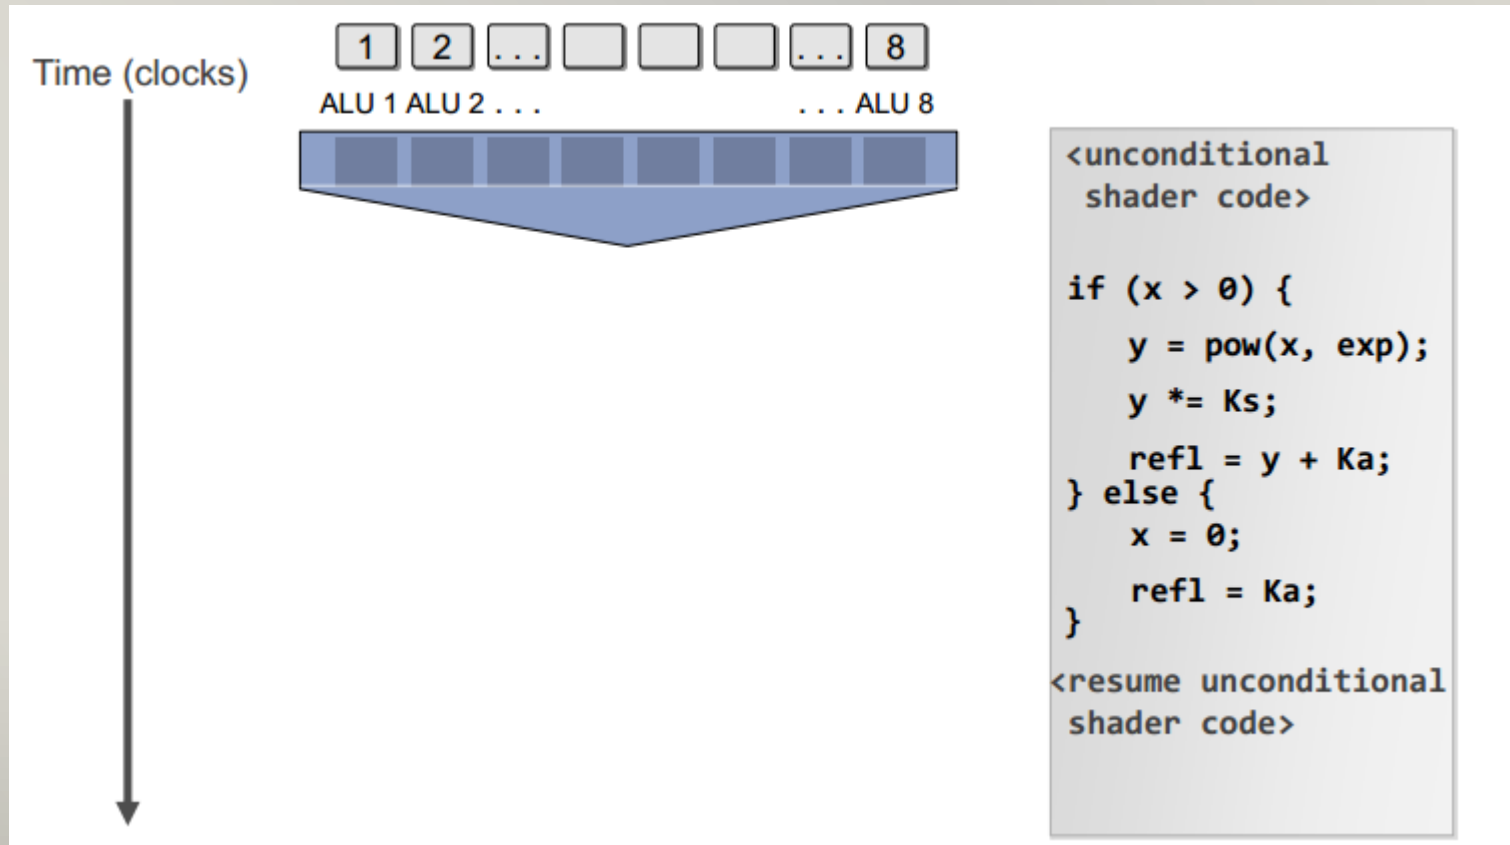
# Cores in GPU

- Each core can process different elements including
  - Vertices
  - Triangles (primitives)
  - Fragments



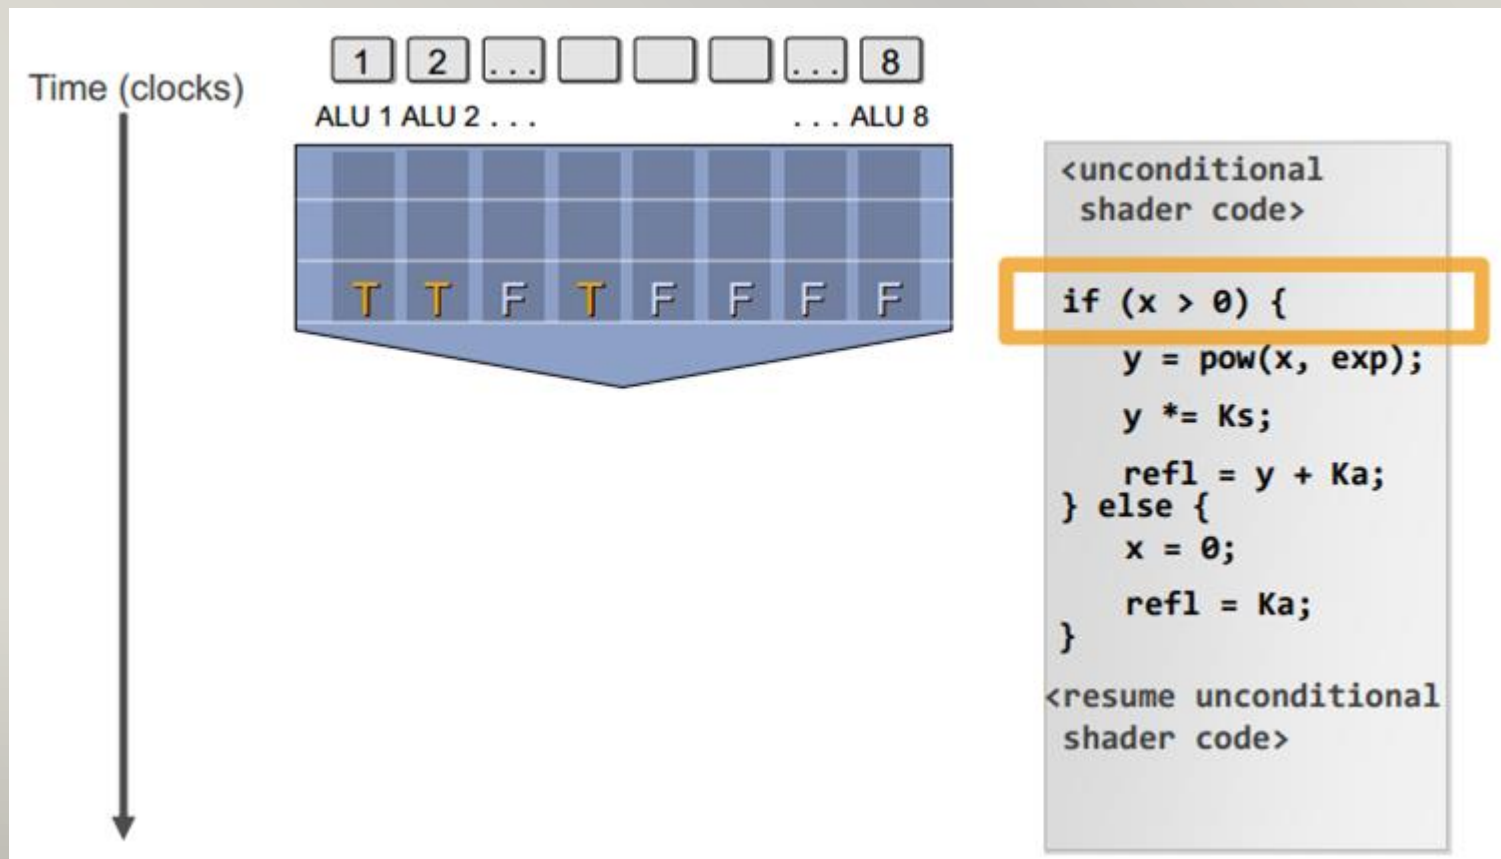
# Problem of Branching in Core

- However, not all element runs the same



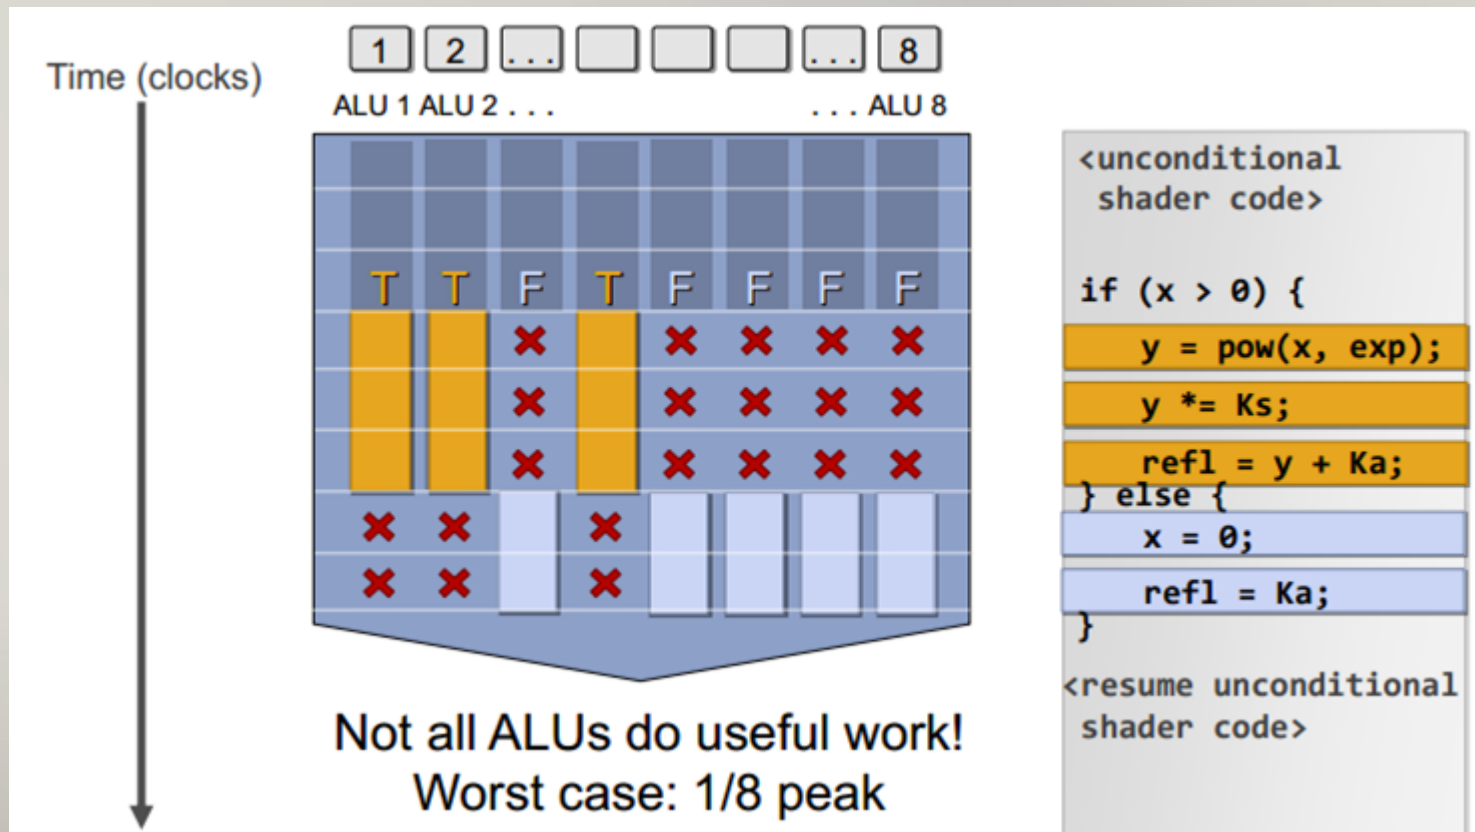
# Problem of Branching in Core

- An if-else will divide elements in 2 groups



# Problem of Branching in Core

- Cross represents ALU idle time





# Solution of Branching in Core

- Increase throughput with enlarged context
- Avoid latency stalls by interleaving execution of many groups of fragments
  - When one group stalls, work on another group

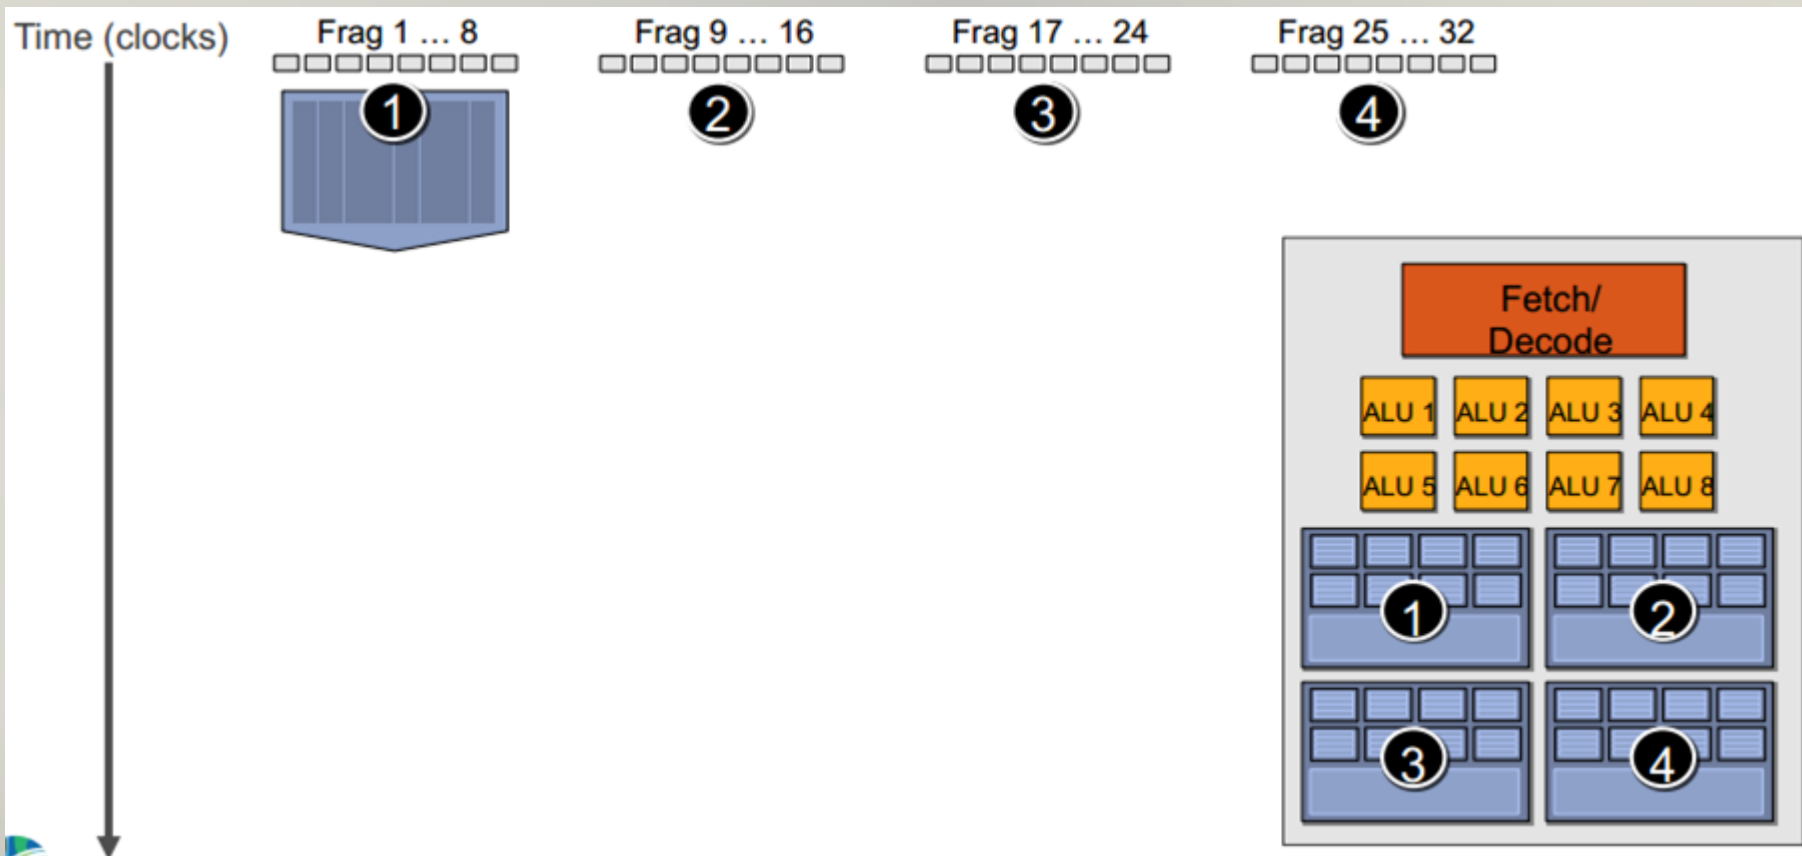


# Summary: three key ideas

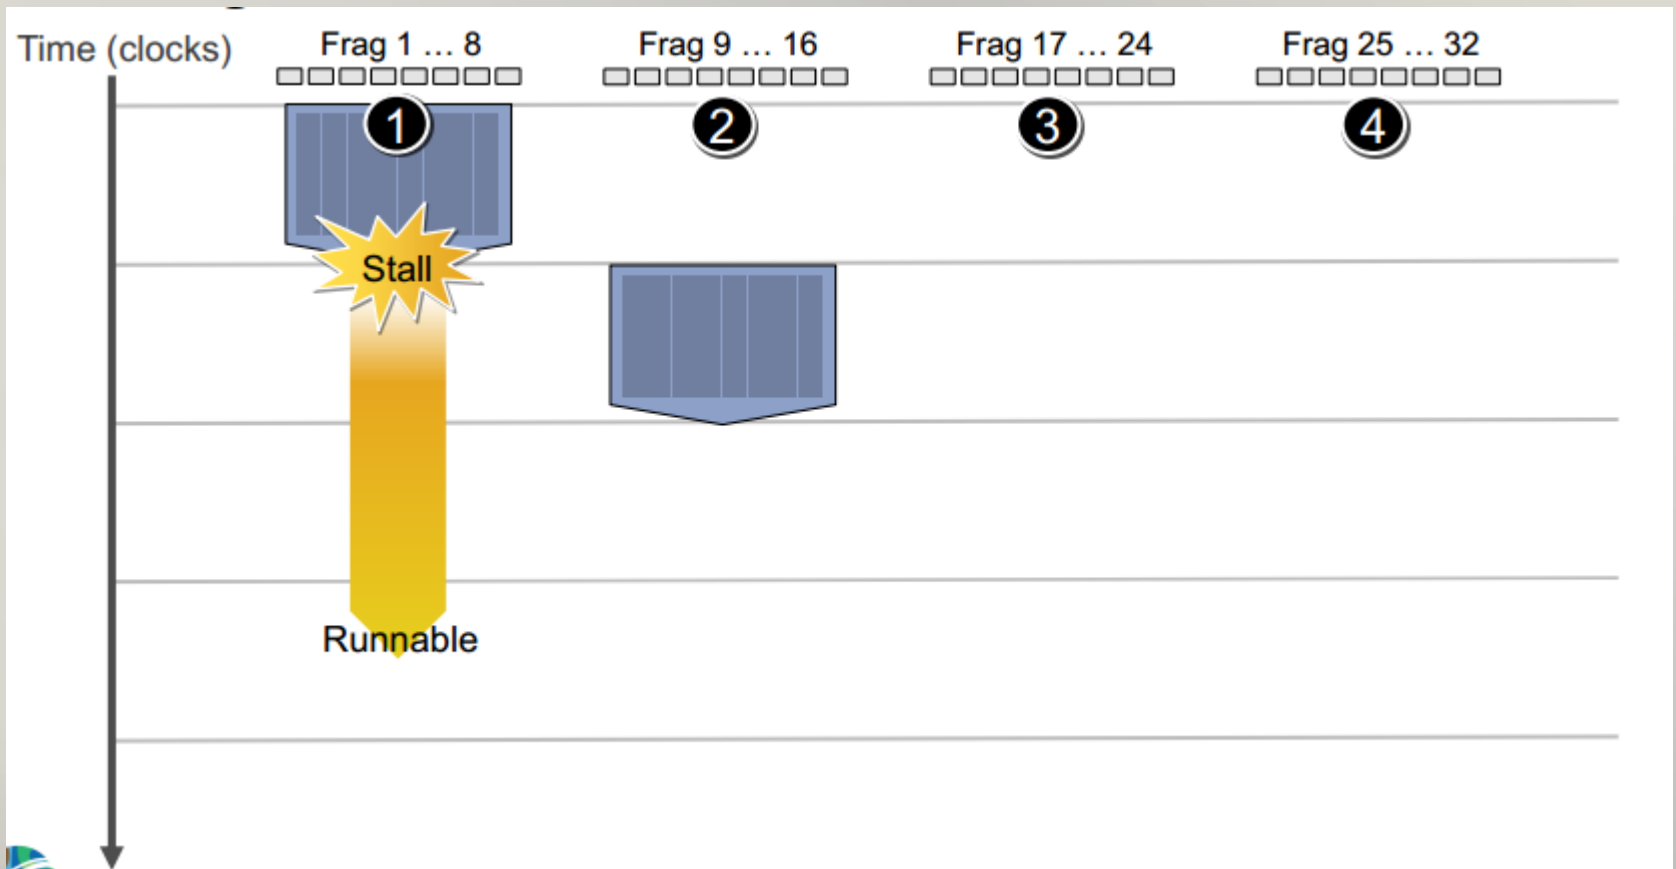
1. Use many “slimmed down cores” to run in parallel
2. Pack cores full of ALUs (by sharing instruction stream across groups of fragments)
  - Option 1: Explicit SIMD vector instructions
  - Option 2: Implicit sharing managed by hardware
3. Avoid latency stalls by interleaving execution of many groups of fragments
  - When one group stalls, work on another group



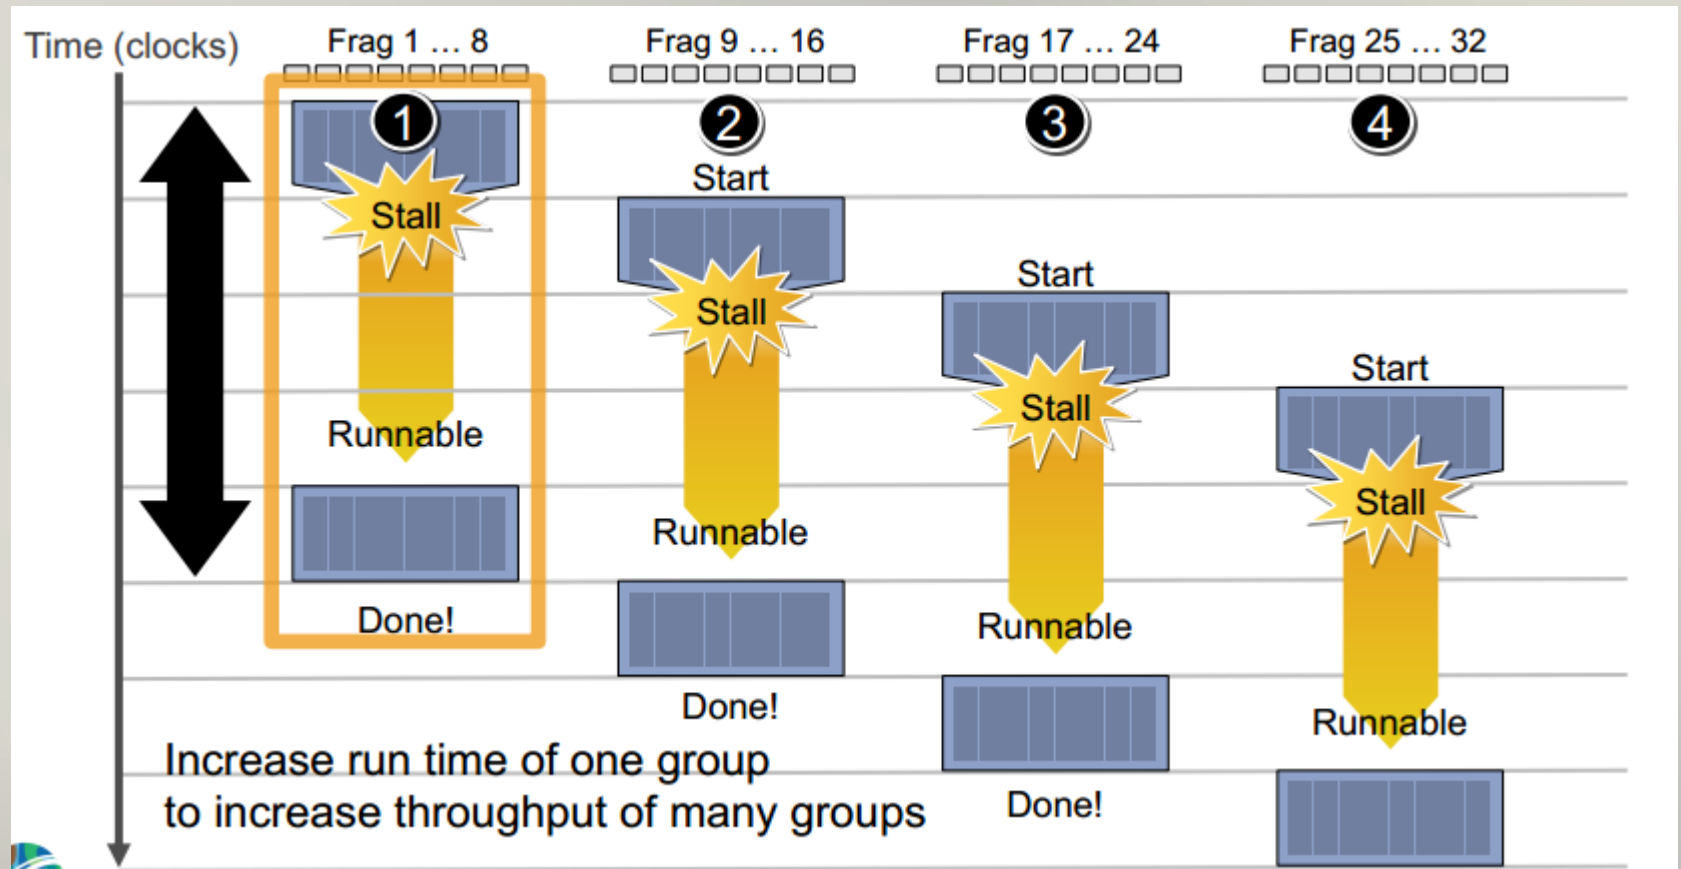
# Solution of Branching in Core



# Solution of Branching in Core

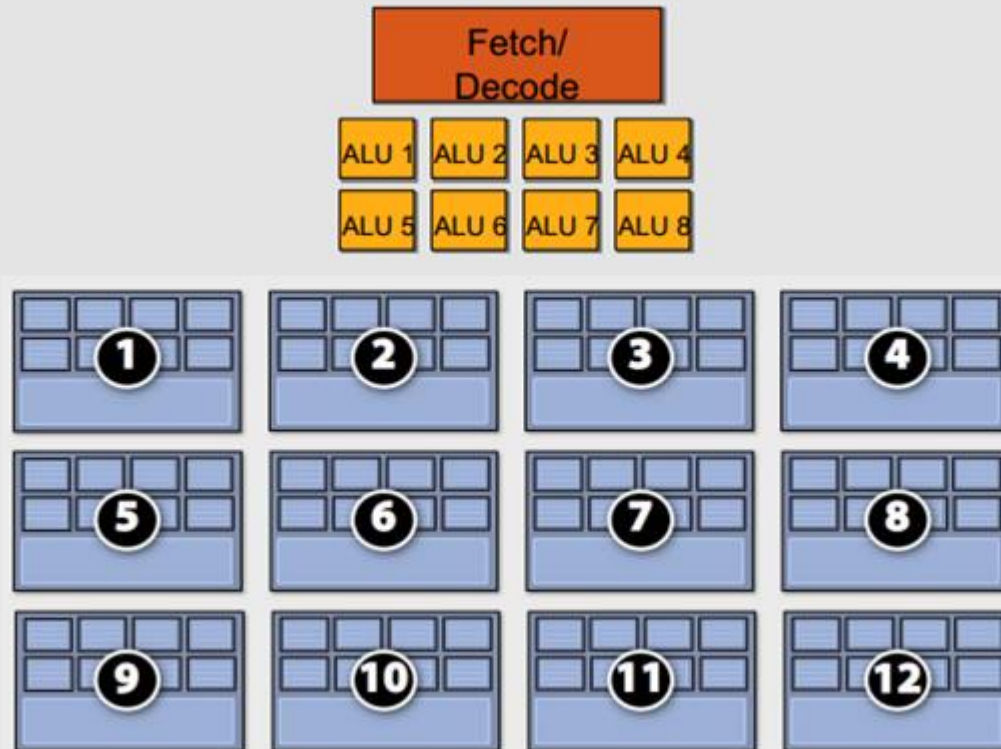


# Solution of Branching in Core



# Solution of Branching in Core

- Storing context



# An Example Chip

16 cores


8 mul-add ALUs per core  
(128 total)

16 simultaneous  
instruction streams

64 concurrent (but interleaved)  
instruction streams

512 concurrent fragments

= 256 GFLOPs (@ 1GHz)







# Shader Programming

- Setup of environment in the rendering engine (e.g. OpenGL)
  - Loading geometry, texture and other graphics elements into the engine (i.e. memory on GPU)
  - Loading shader program
  - Bind the shader program in rendering loop
- Shader program runs in parallel for all elements automatically during rendering
- The typical output is to frame buffer, but for GPGPU, we may have to dump output back to main memory





# Shaders

- GLSL is used as example in the follow
- *Shader object*: an individual vertex, fragment, etc. shader
  - Are provided shader source code as a string
  - Are compiled
- Shader program: Multiple shader objects linked together
- Uniform variables: for passing in parameters to the shader program

# Shader Objects

## ■ Compile a shader object:

```
const char *source = // ...  
GLint sourceLength = // ...
```

Create a Shader Object and return its handle

```
GLuint v = glCreateShader(GL_VERTEX_SHADER);
```

```
glShaderSource(v, 1, &source, &sourceLength);
```

Provide the shader's source code as a string

```
glCompileShader(v);
```

```
// ...
```

```
glDeleteShader(v);
```

Compile the shader

# Shader Programs

## ■ Link a shader program:

```
GLuint v = glCreateShader(GL_VERTEX_SHADER);  
GLuint f = glCreateShader(GL_FRAGMENT_SHADER);  
// ...
```

```
GLuint p = glCreateProgram();  
glAttachShader(p, v);  
glAttachShader(p, f);
```

A program needs a vertex and fragment shader

```
glLinkProgram(p);  
// ...  
glDeleteProgram(v);
```

Link the shaders to form a shader program

**Notice that the program had not yet applied to any of your rendered objects !!!**

# Using Shader Programs

Apply the Shader Program before any draw commands of OpenGL.

The same program can be applied on different drawing targets

```
GLuint p = glCreateProgram();  
// ...
```

```
glUseProgram(p);
```

```
glDraw*(); // * because there are lots of draw functions  
// ...
```

# Uniforms

```
GLuint p = glCreateProgram();  
// ...  
glLinkProgram(p);
```

Each *active* uniform variable has an integer index location (or handle).

```
GLuint m = glGetUniformLocation(p, "u_modelViewMatrix");  
GLuint l = glGetUniformLocation(p, "u_lightMap");
```

```
// ...
```

```
glUseProgram(p);  
mat4 matrix = // ...
```

glUniform\* for all sorts of datatypes

```
glUniformMatrix4fv(m, 1, GL_FALSE, &matrix[0][0]);  
glUniform1i(l, 0);
```

**Uniforms can be changed as often as needed before applying to drawing commands, but are constant during a draw call !!**

# Shader Programming

## A Simple Example

- Input: a teapot model
- Vertex Shader: Manipulate the vertices of the model so as to make it flatten
- Fragment Shader: a pass through or assigned certain color, e.g.



```
void main() {  
    gl_FragColor = vec4(0.4, 0.4, 0.8, 1.0);  
}
```



# The Flatten Shader

- The Vertex Shader makes all input vertices with  $z = 0$  (keeps  $x$  and  $y$ )
- Actually making a cross-section at  $z=0$



<http://www.lighthouse3d.com/>

```
void main(void) {  
    vec4 v = vec4(gl_Vertex);  
    v.z = 0.0;  
    gl_Position = gl_ModelViewProjectionMatrix * v;  
}
```

Shader predefined values are started with "gl\_"

Turn every coordinate with  $z = 0$

# The Flatten Shader

- Another alternative is to assign z value with a sine function of coordinate x
- It will produce a wavy flat teapot



```
void main(void) {  
  
    vec4 v = vec4(gl_Vertex);  
    v.z = sin(5.0*v.x)*0.25;  
  
    gl_Position = gl_ModelViewProjectionMatrix * v;  
  
}
```

# Simple Toon Shading

- Input: a teapot model
- Vertex Shader: Multiply vertex with modelview and projection matrices, and pass normal to fragment shader
- Fragment Shader: Apply a stepping function to map the rendered color



# Simple Toon Shading

## ■ The Vertex Shader

- Passing normal vector to fragment shader
- Values will be interpolated on each fragment

```
varying vec3 normal;  
void main() {  
  
    vec4 v = vec4(gl_Vertex);  
    gl_Position = gl_ModelViewProjectionMatrix * v;  
  
    normal = gl_Normal;  
}
```

# Simple Toon Shading

- Based on per-fragment normal to compute shading
- Then, use shading intensity to map final color

```
uniform vec3 lightDir;  
varying vec3 normal;  
void main() {  
    float intensity;  
    vec4 color;
```

Compute shading using Phong model

**`intensity = dot(lightDir, normalize(normal));`**


```
if (intensity > 0.95) color = vec4(1.0, 0.5, 0.5, 1.0);  
else  
if (intensity > 0.5) color = vec4(0.6, 0.3, 0.3, 1.0);  
else  
if (intensity > 0.25) color = vec4(0.4, 0.2, 0.2, 1.0);  
else  
    color = vec4(0.2, 0.1, 0.1, 1.0);
```

```
gl_FragColor = color; }
```

Smoothly changing  
intensity is mapped to 4  
different colors

# Simple Toon Shading

- The mapping is like a **quantization** to the shading intensity





# GPGPU

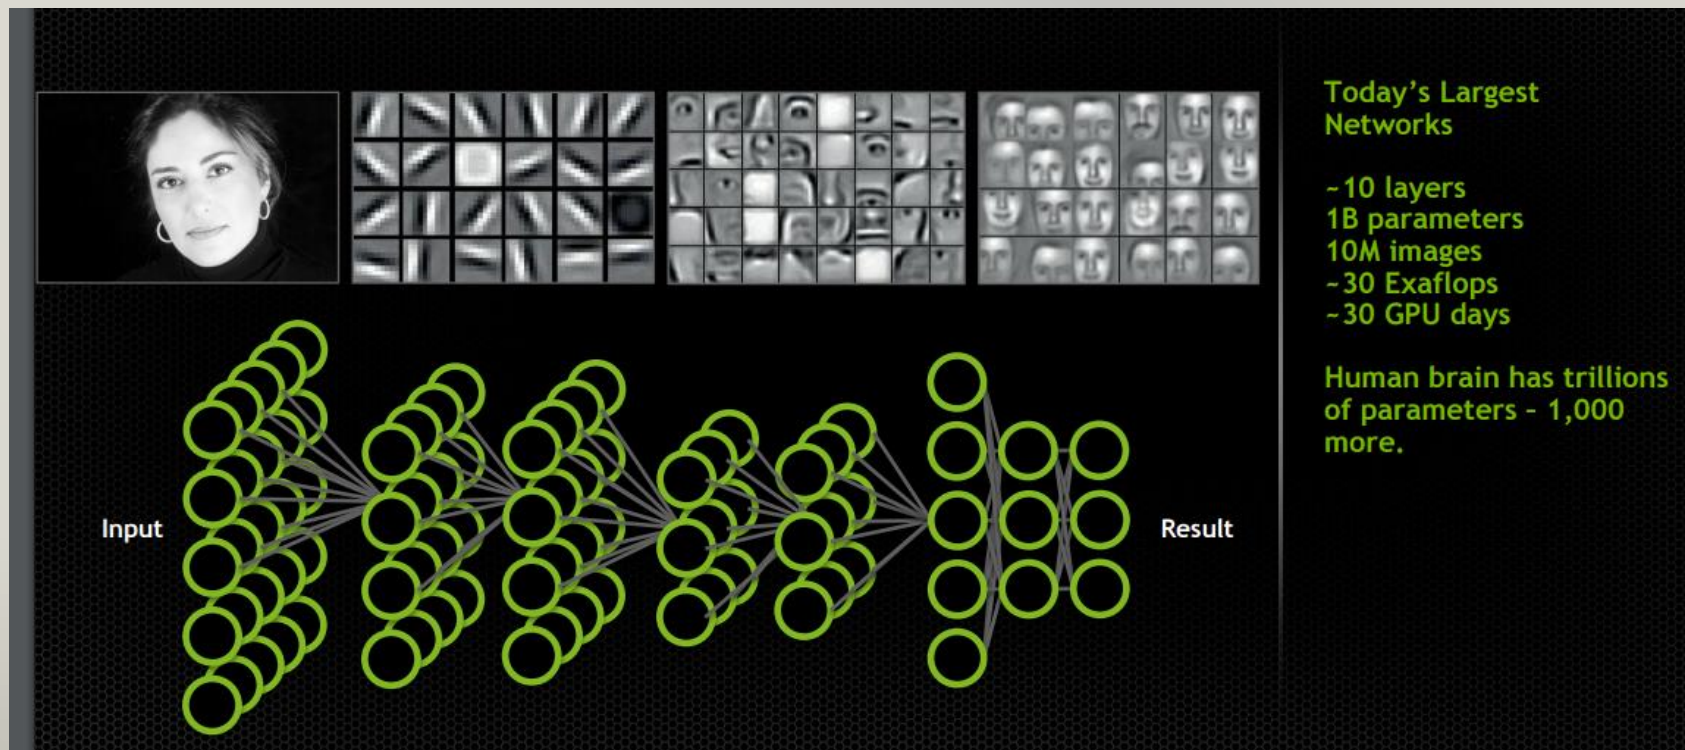
- General Purpose GPU
  - Not only for graphics related applications
  - Physics based simulation can also be regarded as a kind of GPGPU
- Recent examples: Bitcoin Miner using GPU



The advertisement is set against a solid orange background. At the top center is a gold Bitcoin icon. Below it, the text 'GUIMiner' is displayed in a large, bold, white font. Underneath 'GUIMiner' is the subtitle 'a GPU/CPU Bitcoin miner for Windows based on *podbm*' in a smaller white font. Below that, the text 'Last version: 20121203 - Download' is shown, with '20121203' in red and 'Download' in blue. To the right of this text is a small rectangular advertisement for Plus500. This ad has a blue background and features the text 'Trade Bitcoin' in white, followed by two bullet points: 'Buy or short' and 'Trade with leverage'. Below these is a green button that says 'Get 25€ signup bonus'. At the bottom of the Plus500 ad, the logo 'Plus500' is visible, along with the website 'www.Plus500.com' and the text 'CFD Service Your capital is at risk'.

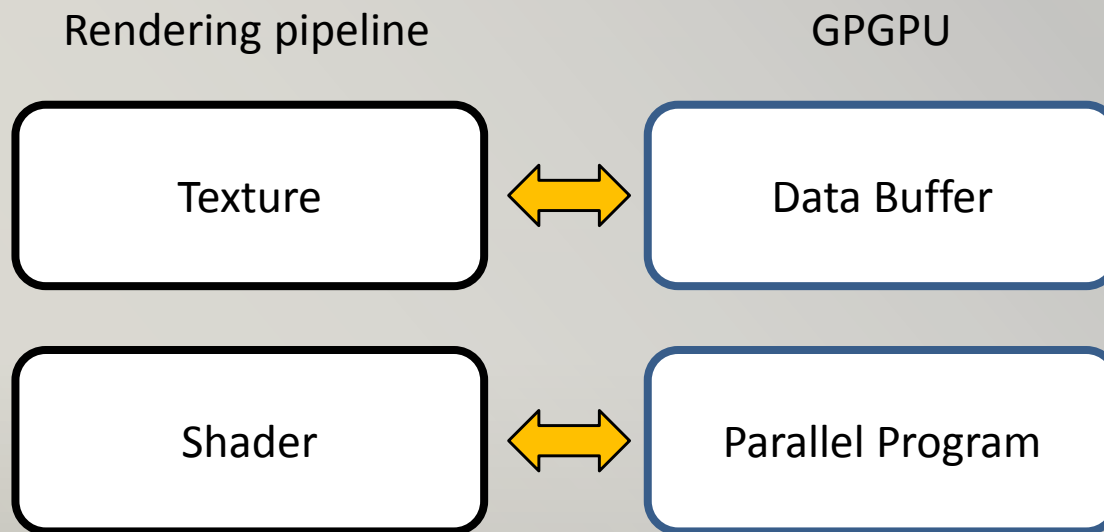
# GPGPU

- Recent examples: Deep learning




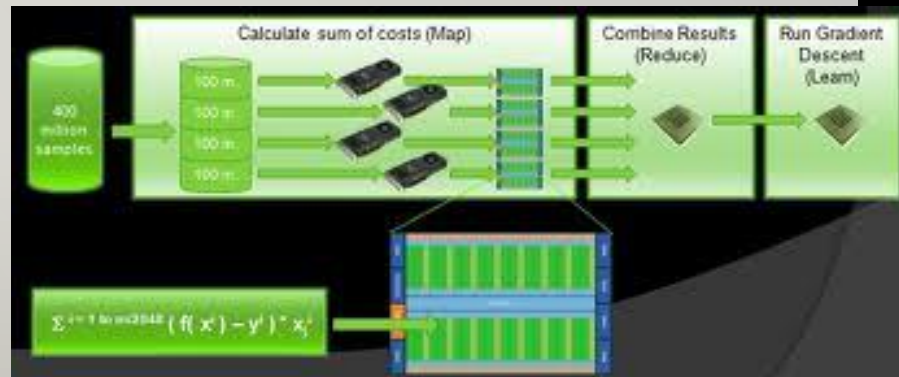
# GPGPU

- We use GPU as a parallel processor
- A mapping between rendering pipeline and general purpose computation
  - Use Texture as Data Buffer, Shader as Parallel Program




# Application of GPGPU

- Accelerated Linear Algebra Computation
  - E.g. CUBLAS, cuSPARSE
- Massive Random Number Generation
- Machine Learning Algorithms
- Sorting




# Pseudo-random number generator (PRNG)


- Provide uniform random numbers
- Example : `rand()` in C
- Stochastic rendering algorithms
  - Photon-mapping
  - Distribution ray-tracing
- Poor randomness -> slow convergence



# Pseudo-random number generator (PRNG)



Results from an LCG  
10,000 photons



Control Image  
10,000 photons



# Some common PRNG

- linear congruential generator (LCG)
  - $R_{n+1} = aR_n + b \pmod{m}$
- lagged Fibonacci generator
  - $R_n = R_{n-j} \# R_{n+k} \pmod{m}$  (where  $\#$  is a binary operator)
- High precision integer arithmetics
- Can't fit in the current GPU !
- GPU implementation is not available




# PRNG on GPU


- CA-based PRNG
- No high precision integer arithmetics
- Adoptable PRNG for different GPU
- PRNG Performance on different GPU
  - 7800 GTX and FX5900
- Why not find optimal PRNG automatically ?

# GPU Accelerated PRNG

- GPU : NO High precision integer arithmetic
- Most PRNG cannot fit in all GPU



Cell D: 1




Cell C: 0




Step(1,  $3 - 1 - 2 \cdot 0$ )

A


# Convergence of Optimization




Control  
10,000 photons




Generation 1




Generation 2



Generation 4



Generation 8



Generation 11

# Performance

- Performance compare with CPU
- 1,000 Parallel PRNG
- 13 times faster

Random numbers generated	GPU CA-PRNG	Software CA-PRNG
10,000	0.004s	0.043s
100,000	0.031s	0.425s
1,000,000	0.31s	4.274s
10,000,000	3.098s	43.003s
100,000,000	31.875s	430s

“Implementing High-Quality PRNG on GPU,”

W. M. Pang, T. T. Wong and P. A. Heng,

***Shader X5: Advanced Rendering Techniques***, Edited by W. Engel, Charles River Media, 2007, pp. 579-590.



# Quicksort algorithm

- Divide and conquer approach (recursive)
  - Recursively divide the sequence into two based on a selected pivot value
- For  $n$  items, it has  $O(n \log(n))$  complexity on average and  $O(n^2)$  in the worst case.
- But it is not very well fit for parallel processing
  - The two sequence are not balanced
  - Operation depends on the previous operations



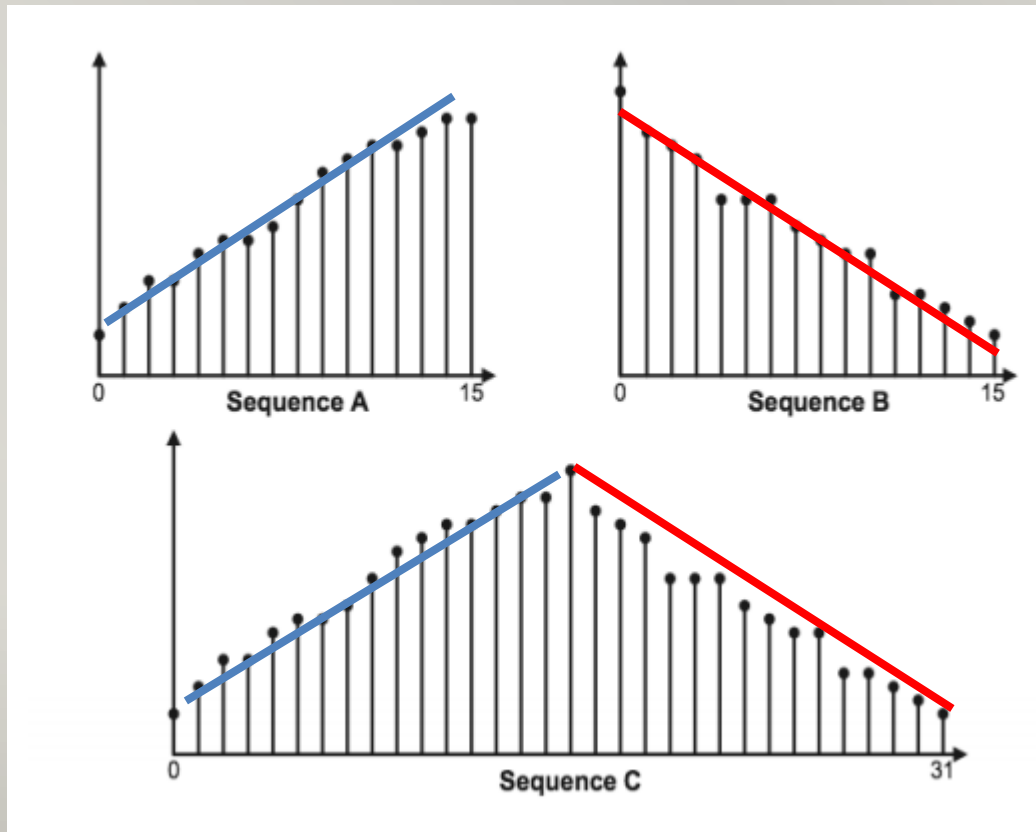


# Bitonic Sort

- Bitonic sort is a sort algorithm that works only with bitonic sequences of values
- A bitonic sequence is
  - A sequence  $a_0, a_1, \dots, a_n$  is bitonic if and only if
  - There is an  $i$ , such that  $a_0, \dots, a_i$  is monotonically increasing and
  - $a_i, \dots, a_n$  is monotonically decreasing
- There is a cyclic shift of the sequence which makes the previous condition true


# Bitonic Sort

- Bitonic sequences




# Bitonic Split

- This operation contains 2 Steps
  - Compare each element in the lower half of the sequence ( $0 - i$ ) with the corresponding element in the upper half ( $i + N/2$ )
  - If the element in the lower half is greater than the element in the upper half, swap the two elements



# Bitonic Split

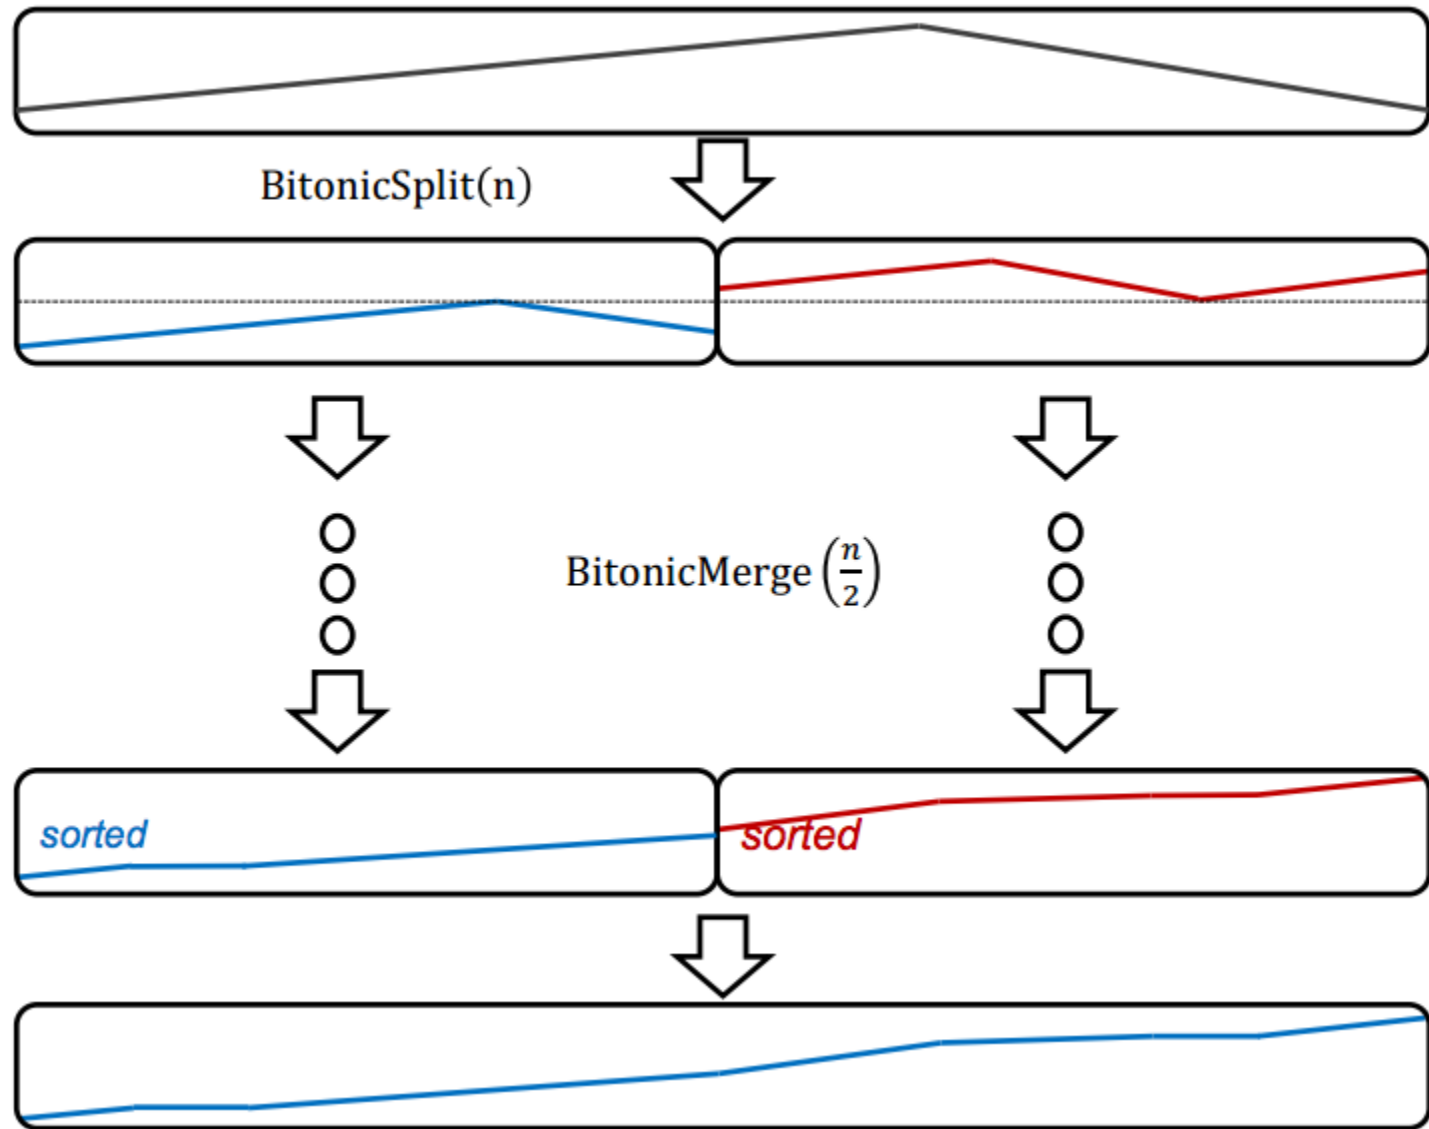
- If we apply the bitonic split to the sequence C , we obtain a sequence which is not bitonic
- But the two sequences are bitonic themselves





# Bitonic Merge

- The bitonic merge consists in applying the bitonic split iteratively
- If a sequence contains  $N$  elements  $N=2^k$ , the bitonic merge requires  $k$  steps
  - Each step is a bitonic split
  - Step  $i$ , splits  $N/i$  elements



The whole sequence will be sorted after the bitonic merge operation !!





# Bitonic Merge

- If the sequence is bigger than  $N=2^k$  but smaller than  $N=2^{k+1}$ 
  - We sort  $N=2^{k+1}$  elements, filling the extra values with the maximum possible value
- It is obvious that the algorithm can be parallelized easily
  - Each comparison in a pair of numbers within bitonic split is independent to each other





# Bitonic Sort

- But what if our sequence is not bitonic at the beginning? It is very likely to be...
- How about we first transform the sequence into a bitonic one?
- Split the sequence into two half each time
- Apply recursively on each half to construct a bitonic sequence

# Bitonic Sort

- Ultimately, we will split into a sequence of only two elements
  - Note that all two-elements sequence is already bitonic
- So, in each step, our objective is to
  - Combine two sequences with the first half being increasing and the upper half decreasing




Our Goal




# Bitonic Sort

- In other words, we are doing two sortings
  1. First half : sort in ascending order
  2. Second half: sort in descending order
- So, we can apply a Bitonic sort (or bitonic merge) !
  - As each half now is bitonic itself




split(n) ↓



BitonicSort( $\frac{n}{2}$ )



BitonicMerge(n) ↓






# Bitonic Sort

- Creating a bitonic sequence is actually sorting the data
  - Following a special order
- It is data independent
  - We can make a sorting network!
- So, very suitable for parallel processing

# Deep Learning with GPU

- cuDNN
  - <https://developer.nvidia.com/cudnn>
- Many deep learning libraries support GPU
  - Caffe, Torch, Chainer
- DIGITS
  - For visualization
- (refers to notes from nVidia)







# Summary

- We have discussed different topics related to GPU, e.g.
  - Difference between CPU and GPU
  - Strength and weakness of GPU
  - Graphics applications of GPU
  - Various types of Shaders
  - General Purpose GPU